

# CS 33 Final Review Guide

Note: This review guide may not contain all topics needed on the exam. It is intended to simply help guide you as you study for your exam.

## I. Bits and Bytes

### a. Representation of Numbers

- i. Can you convert numbers between binary, decimal, and hexadecimal forms?

### b. How computers represent data types of C (particularly for the x86-64 architecture)

### c. Data Operations

#### i. Boolean Operations

1. And (&)
2. Or (|)
3. Xor (^)
4. Not (~)
5. Arithmetic and Logical Shifts (<<, >>)
  - a. For which types of data would you do arithmetic shifts and for which would you do logical shifts?

#### ii. Logical Operations

1. And (&&)
2. Or (||)
3. Not (!)

#### iii. What is the difference in the output of boolean operations and logical operations?

#### iv. Memory and Strings

1. How is memory addressed (particularly for x86-64)?

## **2. Endianness**

**a. Little Endian vs. Big Endian**

**b. What Endianness is used for x86-64?**

## **3. Strings**

**a. How are strings are denoted and represented in C?**

# **II. Integers**

## **a. Unsigned**

**i. Minimum and Maximum represented numbers for an n bit integer.**

## **b. Signed**

**i. 2's complement representation of an integer.**

**1. How to calculate the negative of a 2's complement number**

**2. Minimum and maximum represented numbers for an n bit 2's complement integer.**

**c. Casting between unsigned and signed integers (implicit and explicit)**

**i. Results when comparing unsigned and signed integers (any combination of the 2)**

## **d. Integer addition**

**i. How is addition performed?**

## **ii. Overflow**

**1. Unsigned overflow vs. Signed Overflow**

**2. How to detect overflow?**

## **e. Integer shifting**

**i. Shifting left by n bits is equivalent to what arithmetic operation?**

**ii. Shifting right?**

# **III. Assembly Basics (x86-64)**

- a. Concepts of registers, PC, and memory.
  - i. What exactly do we store in registers?
- b. Assembly instructions
  - i. Basic format
    - 1. Where src and dest values are located.
    - 2. Instruction suffixes:
      - a. What do the b, w, l, q suffixes for instructions mean?
      - b. Can you identify which portion of a register is being used by its name?
        - i. (i.e. %rax vs. %eax)
  - ii. lea vs. mov
  - iii. Be familiar with other basic operations like add, sub, etc.
  - iv. Calculating memory addresses
    - 1. When in assembly would you dereference an address?
    - 2. Format and formula for calculating memory addresses.
      - a. Ex. What is the address calculated by `20(%rax, %rcx, 2)`?

#### **IV. Control**

- a. Condition Codes
  - i. CF, SF, ZF, OF
    - 1. What are these control codes and when are they set?
    - 2. Explicit and Implicit code setting
- b. Control Instructions
  - i. Instructions that explicitly set condition codes
    - 1. test, cmp

- a. What operations do these instructions perform, and how do they set the condition codes?
      - b. Do these instructions have any side effects besides setting condition codes?
    - ii. Instructions which use condition codes.
      - 1. Jump instructions (jX)
        - a. e.g. jg, jle, jne, je, etc.
      - 2. Other conditional instructions
        - a. SetX
        - b. Conditional Move
          - i. Ex. Cmovle
      - 3. How can you use these instructions to implement branches, loops and switch statements?

## **V. Procedures**

- a. Stack
  - i. Stack organization
    - 1. Direction stack grows and location of the top of the stack.
    - 2. %rsp register
      - a. What does it contain?
  - ii. Stack commands
    - 1. pushq
    - 2. popq
    - 3. What does using sub on %rsp do?
- b. Procedure control
  - i. Instructions
    - 1. callq
    - 2. retq

- 3.** What do these instructions do? How do they affect the PC (%rip), %rsp, and the stack?
- ii.** Data conventions
  - 1.** In which registers are procedure arguments stored in?
  - 2.** In which register is the return value stored in?
- iii.** Be familiar with the overall process of passing control when a procedure is called and when it is returned.
- iv.** Caller-saved registers vs. Callee-saved registers.
- v.** Concept of stack frames
  - 1.** What type of data is typically stored in a stack frame?
- vi.** Recursion
  - 1.** How recursion looks as assembly code and how it is represented as stack frames?

## VI. Data

- a. Arrays**
  - i.** Is data contiguous?
  - ii.** How do you measure size of array?
  - iii.** How is array stored in memory?
  - iv.** C syntax
    - 1.** Suppose you have array: `int val[5]`
      - a.** Indexing: what does `val[3]` return?
      - b.** What does `val` return?
      - c.** What does `val + 2` return?
      - d.** What does `*(val + 1)` return?
      - e.** What does `&(val[0])` return?
  - v.** Multidimensional arrays:
    - 1.** How are multidimensional arrays stored in memory?

- a. What flattening method is used?
- 2. If given an index for a value or address in an array, can you find it in memory?

**b. Structs**

- i. How is data stored in a struct?
  - 1. Is data contiguous?
  - 2. Memory representation
    - a. Alignment (Which addresses would you store the data?)
      - i. What are the padding rules for individual elements?
      - ii. What are the padding rules for the overall struct?
    - b. How would you access an element if you only had a pointer to the struct?
    - c. Reorder elements to save space.

**c. Unions**

- i. How is data stored in unions?
  - 1. Which address(s) would you store the data?
- ii. How to measure size of unions?

**VII. Advanced Machine Programming Topics**

**a. Memory space.**

- i. How is memory organized?
  - 1. Stack
  - 2. Heap
  - 3. Data
  - 4. Text/Shared Libraries
- ii. Which types of data goes into which location?
  - 1. I.e., where would local procedure data be stored?
  - 2. Where would static variables be stored?

3. Where would procedures and instructions be stored?

**b. Buffer Overflow**


- i. How does buffer overflow manifest?
- ii. What are some issues with buffer overflow?
  - 1. i.e. Stack smashing
- iii. What are some defenses for buffer overflow?

**VIII. Floating Point**

**a. IEEE Floating Point Formula**

- i. Normalized Floating Point
- ii. Denormalized Floating Point
- iii. Special Cases (Inf, -Inf, and NaN)
- iv. In each case, what is the value of the exponent?
- v. For Normalized and Denormalized Floating Point, What is the leading value of M?
- vi. Given an exponent field size, how do you calculate bias and E in the floating point formula?
- vii. How do you determine the sign of a floating point value?

**b. Useful Visualization**



**c. Properties of Floating Point operations**

- i. Rounding
  - 1. Rules

- a. Towards 0
- b. Round down
- c. Round up
- d. Nearest Even

## 2. Rounding in binary

### ii. Addition

- 1. How to add up the M fields in conjunction with the E fields, and round M to remove excess bits.

### iii. Multiplication

- 1. Basic understanding of what is done.

### iv. How are Inf and NaN values generated from addition and/or multiplication ops?

### v. Casting

- 1. Cast from float to double and vice versa?
- 2. Cast from float/double to int and vice versa?
- 3. In which cases will the value be preserved?
- 4. Rounding and truncation rules when casting in general.

### d. Comparisons and Equations

- i. What is the result of a comparison with Inf?
- ii. What is the result of a comparison with NaN?
- iii. Can you determine if 2 expressions are equivalent?
  - 1. Study examples from lecture and discussion.

## IX. Optimization

### a. General Optimizations

- i. Code/Motion
- ii. Strength reduction
- iii. Sharing common subexpressions
- iv. Removing unnecessary procedure calls

### b. Optimization Blockers



- i. Procedure Calls
  - 1. Why are procedure calls expensive?
  - 2. Why can't compilers optimize these easily?
- ii. Memory Aliasing
  - 1. When multiple references refer to same or overlapping memory?
  - 2. Why can't compiler optimize away certain memory accesses if the programmer knows they don't refer to the same memory?
  - 3. Look at examples.
- c. Instruction Level Parallelism
  - i. In general, why can we do this?
  - ii. Cycles Per Element/ Cycles per op
    - 1. In general:  $T$  (Total execution time) =  $CPE * n(\text{number of iterations}) + \text{overhead}$
    - 2. Would adding  $m$  functional units improve  $T$  by a factor of  $m$ ?
  - iii. Latency and Throughput
    - 1. Latency is thought of as how long a function/operation takes overall.
    - 2. Throughput is how many ops can be completed in a time unit (i.e. cycle).
    - 3. If there was no parallelism,  $\text{Throughput} = 1/\text{Latency}$ .
    - 4. How about if there was parallelism?
  - iv. ILP optimizations
    - 1. Loop-unrolling
      - a. What is it?
      - b. Why does this improve performance?

- i. Does it remove sequential dependencies?
- c. Does this improve latency and/or throughput?

## 2. Reassociation

- a. What is it?
- b. Why does this improve performance?
  - i. Does it remove sequential dependencies?
- c. Does this improve latency and/or throughput?
- d. Consider cases of simple reassociation of operations vs. using multiple accumulators.

## 3. Estimate improvement of performance using latency and throughput bounds (given).

- a. Understand how/why optimizations may improve/reduce performance.

## v. Additional Topics

### 1. SIMD

- a. Vector operations
- b. How does this improve performance?

### 2. Branch

- a. Why are these operations typically so expensive?
- b. Branch Prediction: How does a CPU use this to improve branch performance?

## X. Memory Hierarchy (Caching in Particular)

### a. Locality

#### i. Temporal

- 1. What does this mean?

**ii. Spatial**

**1. What does this mean?**

**iii. Consider examples for both data and instruction memory.**

**b. Hierarchy**

**i. Understand at a high level how the cache hierarchy works.**

**1. How is the hierarchy organized?**

**2. Why does the hierarchy work?**

**ii. Memory hits vs. misses**

**1. Understand at a high-level different policies for hits and misses exists, which includes how blocks are evicted from the cache and replaced.**

**c. Cache Organization**

**i. Cache Structure**

**1. What is a block/line?**

**2. What is a set?**

**ii. Cache addressing and accessing**

**1. Which bits of a memory address determine the block offset?**

**a. Given a block size  $B$ , how many offset bits do you need at least?**

**2. Which bits of a memory address determine which set you are accessing in the cache?**

**a. How is this done?**

**b. If you have  $S$  sets, how many set bits do you need at least?**

**3. Once you know which set the address is referring to, how do you determine which block the set is referring to?**

- a. What happens if the block is not found in the set?
  - 4. What is the formula relating cache size  $C$  to block size  $B$ , set size  $E$ , and number of sets  $S$ ?
    - a. Given certain values about the cache structure and memory address structure, calculate the remaining ungiven values.
  - 5. Familiarize yourself with cache organization on multi-core system
    - a. What is False Sharing?
- d. Practical Optimization using Cache Structure
  - i. Given code example, modify to take advantage of cache locality.
  - ii. Particular examples for multi-dimensional arrays.
  - iii. Estimate misses of certain implementation given block size.

## **XI. Multithreading**

- a. High-level overview
  - i. Understand how to decompose program into multiple threads
    - 1. Domain decomposition
    - 2. Task decomposition
    - 3. Pipelining
  - ii. Fork-Join Multithreading Model
    - 1. What does this mean? How does it work?
  - iii. What is a dependency graph? How is it used for decomposition?
- b. OpenMP General
  - i. Know the various pragma's, macros and clauses used by OpenMP for multithreading

**1. Highlights (this is not an exhaustive list of all macros required):**

- a. parallel for**
- b. parallel**
- c. for**
- d. private, firstprivate, lastprivate**
- e. single**
- f. nowait**

**c. Race Conditions and Locking**

**i. What is a race condition?**

**1. Why does it occur?**

**ii. Locking**

**1. Could we use a simple flag variable to prevent race conditions?**

**a. Give a simple logical reason or counterexample.**

**2. If a thread is in a locked section, can another thread enter that locked section?**

**3. Implementing locking with OpenMP**

**a. Macros**

**i. critical**

**1. How does locking work with this macro?**

**2. What is a benefit and downside of this approach?**

**ii. omp\_lock\_t (lock objects)**

**1. How is this different from the critical section?**

**2. What is the benefit of this approach?**

### **3. Downside?**

- iii. Reductions:** How is this similar to the above approaches? Different?

#### **d. Deadlock**

- i. What is deadlock?**
- ii. Why does it occur?**
  - 1. Four conditions**
    - a. Mutual Exclusion**
    - b. Partial Resource Allocation**
    - c. No Pre-empting locks**
    - d. Cycle in Resource Allocation/Locking**
  - 2. Given a code sample, could you identify the deadlock?**
    - a. Prove it using the four deadlock conditions**
    - b. Remove the deadlock**
  - 3. If your code only uses one of the 3 OpenMP locking macros described above, for which of the macros is deadlock possible? Why?**
  - 4. Examples from class/discussion**
    - a. Circular Dining Table/ Circular Spongebob cooking**

## **XII. Linking**

- a. Why use Linkers? What is benefit?**
- b. Static Linking vs. Dynamic Linking?**
  - i. List benefits and costs of both**
    - 1. Think of storage, performance, etc.**
  - ii. Types of library files used for each**
  - iii. At what point of the program does linking occur in each case?**

1. For dynamic linking, there are multiple possible points at which linking can occur. What are they?
  - c. Overview of linking steps
    - i. Symbol Resolution
      1. Types of Symbols
        - a. Global
        - b. External
        - c. Local
      2. What does resolution actually mean?
    - ii. Relocation
      1. What is happening in this step?
      2. What does the program look like before and after the linking?

### **XIII. Exceptions**

- a. What is Exceptional Control Flow?
  - i. How is it different from branching and jumping?
- b. Compare and contrast types of exceptional control flow
  - i. Exceptions
  - ii. Process Context Switch
  - iii. Signals
  - iv. Nonlocal jumps
- c. Exceptions
  - i. What are they?
  - ii. How are they run (handled)?
  - iii. Types:
    1. Interrupts
    2. Traps
    3. Faults
    4. Aborts
  - iv. Compare types of exceptions

1. Asynchronous vs. Synchronous
  2. How they manifest in general?
  3. Return behavior in general
  4. Examples
- d. Understand general examples of exceptions
- i. System Calls
  - ii. Page Faults

#### **XIV. Virtual Memory**

- a. Why do we use it? What do we get from virtual memory?
- b. Benefits and Costs of Virtual Memory
- c. How VM works?
  - i. Page Table
    1. Comparison to caching
      - a. Page offset
      - b. Page Numbers
    2. Page Translation from virtual address to physical one
    3. Page Hits vs. Page Faults
      - a. What happens in each
      - b. Compare to cache hits and misses
    4. How are pages allocated during a page fault?
    5. Permission bits: read, write and execute
    6. TLB and multi-level page table
      - a. What are the purposes of each?
      - b. What benefits do they provide?

#### **XV. RISC Architectures and MIPS**

- a. Compare and contrast RISC/CISC architectures?
  - i. Example: # of available instructions
  - ii. Size of typical programs.
  - iii. Etc.



- b.** Is x86-64 CISC or RISC?
- c.** Is MIPS CISC or RISC?
- d.** MIPS
  - i.** Know the different types of instructions and their formats?
    - 1.** Memory accesses
    - 2.** Arithmetic
    - 3.** Logical
    - 4.** Comparison
    - 5.** Control (branching, jumping)
  - ii.** Can you convert x86-64 code to MIPS and vice versa?
  - iii.** Can you read a MIPS program, understand what it does, and find a solution given some input?
    - 1.** Similar to our x86-64 assembly problems, but with MIPS.

## **XVI. Practical Experience**

- a.** Linux and gdb
  - i.** Basic gdb commands
  - ii.** Can read output of linux and gdb commands, particularly those shown in class.
    - 1.** Can read and interpret register and memory dumps from gdb.
- b.** Data Lab
  - i.** Familiarize yourself with all the functions. You should know the basic algorithm for implementing each of these functions.
- c.** Bomb Lab
  - i.** Be able to understand and interpret functions written in assembly.

- ii. Given input register values, be able to calculate output register values, and vice versa.
  - iii. Particular examples: recursions, switch jump tables, loops, ...
  - iv. Be able to read memory and stack dumps.
- d. Attack Lab
  - i. Buffer Overflow tests
    - 1. Be able to read an assembly algorithm and stack trace, and understand if and how buffer overflow occurs.
    - 2. Be able to recognize a code injection attack: usually if you are jumping into the region of the stack you are overflowing.
    - 3. Be able to recognize a return-oriented programming attack. Follow along the gadget instructions and returned addresses to see what the function is doing.
- e. Parallel Lab
  - i. Why was it so difficult to parallelize the base algorithm given to you?
  - ii. Understand the modification tip given to you to parallelize the code.
    - 1. Why did we do it this way?
  - iii. Understand tradeoffs between parallelization versus cache locality.