



UPE Tutoring:

CS 33 Midterm

Sign-in <https://forms.gle/zzwiqjHCyXgnEPyX6>

Slides link available upon sign-in



Integers and Binary Operations



Integers

- 4 bytes long - 00000000 00000000 00000000 00000000
- Signed vs Unsigned.
- Performing operations between signed and unsigned casts both to unsigned

Signed

- [-2,147,483,648, 2,147,483,647]
- MSB Signed vs Two's complement
 - **MSB** - use most significant (leftmost) bit to denote negativity.
0b100000000000000001 = -1
 - **Two's complement** - invert and add 1

Unsigned

- [0, 4,294,967,295]



Example of Two's Complement

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1

- Two's complement is commonly used. What advantages does it have?



Example of Two's Complement

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Invert

1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

- Two's complement is commonly used. What advantages does it have?



Example of Two's Complement

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Invert

1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Add 1

- Two's complement is commonly used. What advantages does it have?



Example of Two's Complement

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Invert

1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Add 1

- Two's complement is commonly used. What advantages does it have?
- Addition of numbers and their negative give us 0



AND (&)/OR (|)

& (AND)

A	B	
0	0	0
0	1	0
1	0	0
1	1	1

| (OR)

A	B	
0	0	0
0	1	1
1	0	1
1	1	1



XOR (^)

^ (XOR)

A	B	
0	0	0
0	1	1
1	0	1
1	1	0

$0b1011 \wedge 0b0100 = 0b1111$

$0b1010 \wedge 0b0000 = 0b1010$

$0b1010 \wedge 0b1010 = 0b0000$

$0b0011 \wedge 0b1111 = 0b1100$



Bitwise vs Logical NOT

~ (Bitwise NOT)

Flips all the bits

`~0b011100001111 = 0b100011110000`

! Logical NOT

Turns any non-zero number into 0.

Turns zero into 1.

`!0b011100001111 = 0b000000000000`

`!0b000000000000 = 0b000000000001`



Left-Shift (<<)

- Result always equivalent to multiplying by powers of 2
 - Even for negative numbers and with overflowing

```
int num = 0xFF00000F;
```

```
num << 4: 0xF00000F0      num * pow(2, 4): 0xF00000F0
```



Right-shift (>>)

- **Logical Right Shift**

- Does not sign-extend
- Applies on unsigned int
- Like division by powers of 2

```
unsigned num = 0xCF00000F;
```

```
num >> 4: 0x0CF00000
```

```
num / 16: 0x0CF00000
```

- **Arithmetic Right Shift**

- Sign is preserved
- Applies on int
- **Not like** division by powers of 2
 - Rounded towards -Infinity instead of zero

```
int num = 0xCF00000F;
```

```
int num = -7;
```

```
num >> 4: 0xFCF00000
```

```
num >> 1: -4
```

```
num / 16: 0xFCF00001
```

```
num / 2: -3
```



Common Bit Manipulation Patterns

- Negative of a number

$-c = \sim c + 1$

- Mask to 0/1 bit

`!!mask`

- Mask from Sign-bit

`num >> 31`

- Getting a bit

```
bool getBit(int num, int i) {  
    int bit = ((1 << i) & num);  
    return !!bit;  
}
```

- Setting a bit

```
bool setBit(int num, int i) {  
    return ((1 << i) | num);  
}
```



Bit Manipulation Patterns

- Branches

```
if (test)
    output = a;
else
    output = b;
```

```
mask = (!!test << 31) >> 31);
output = (mask & a) | (~mask & b);
```

- Online resource for solutions of harder bit-manipulation problems:
<https://graphics.stanford.edu/~seander/bithacks.html>



Practice Question

Write a C function using only bitwise operators and “!”, to determine if the number given is non-negative. If the number is non-negative return 1 else 0.

Hint: It can be done in one line!



Answer

Write a C function using only bitwise operators and “!”, to determine if the number given is non-negative. If the number is non-negative return 1 else 0.

```
int is_nonnegative(int x) {  
    return !(x >> 31);  
}
```

Let's try something a little harder!



Practice Question

Write a C function using only bitwise operators to return the absolute value of a number. If the number is Tmin, just return Tmin since there is no positive representation of Tmin.

```
int my_abs(int x) {  
    ...  
}
```



Practice Question

Write a C function using only bitwise operators to return the absolute value of a number. If the number is Tmin, just return Tmin since there is no positive representation of Tmin.

```
int my_abs(int x) {  
    int sign = x >> 31;  
    return (x ^ sign) + (1 & sign); // multiply by -1 if the sign is negative  
}
```



Practice Questions

Assuming:

```
int x = rand();
```

```
int y = rand();
```

```
unsigned ux = unsigned(x);
```

```
unsigned uy = unsigned(y);
```

Which of the following are always true?

\Rightarrow means implies

- $(x > 0) \ \&\& \ (y > 0) \Rightarrow x * y > 0$
- $x > uy \Rightarrow x > 0$
- $(x \wedge x) \wedge x == y \Rightarrow x == y$
- $(ux - uy) == x - y$
- $(x \ll 32) \gg 32 == x$



Practice Questions

Assuming:

```
int x = rand();  
int y = rand();  
unsigned ux = unsigned(x);  
unsigned uy = unsigned(y);
```

Which of the following are always true?

=> means implies

- $(x > 0) \ \&\& \ (y > 0) \Rightarrow x * y > 0$ **False (undefined)**
- $x > uy \Rightarrow x > 0$ **False**
- $(x \wedge x) \wedge x == y \Rightarrow x == y$ **True**
- $(ux - uy) == x - y$ **True**
- $(x \ll 32) \gg 32 == x$ **Undefined**



ISAs/x86-64



ISA - x86-64

- ISA - Instruction Set Architecture. Interface between software and hardware
- “CISC” architecture - relatively complex instructions. Opposite would be RISC
- Little-Endian
- Used by Intel Chips
- 64-bit



Byte Ordering

Big Endian

- Least significant byte has highest address

Little Endian

- Least significant byte has lowest address

Example - 0x01234567 at 0x100

0x100	0x108	0x110	0x118
01	23	45	67
67	45	23	01

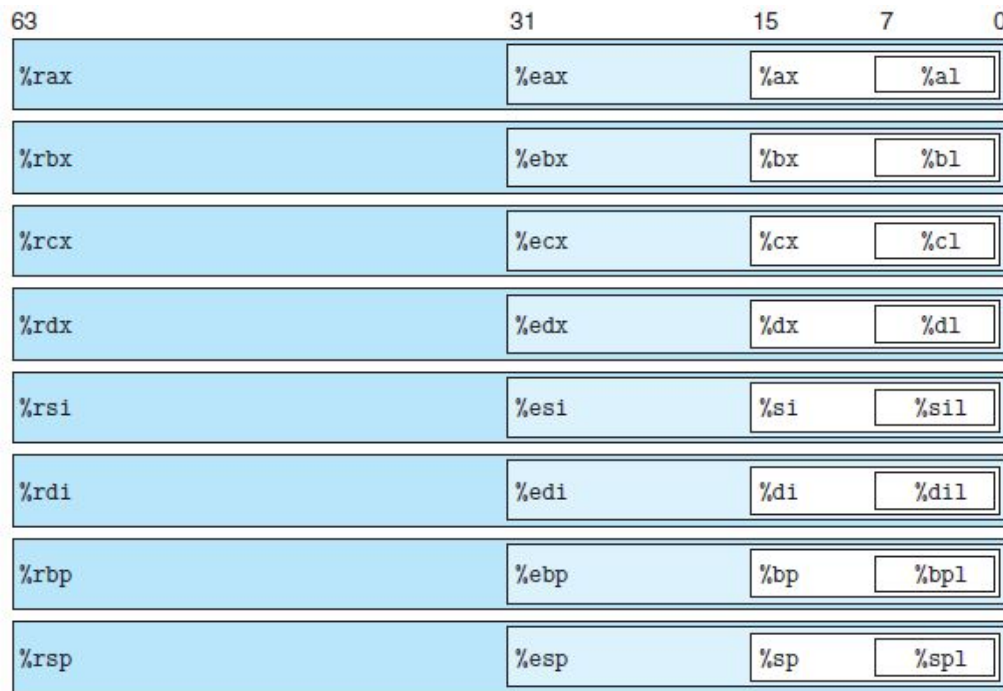
Big Endian

Little Endian



Registers

- Registers are small hardware based storage devices closest to the CPU that generally store addresses into physical memory.
- Assembly code suffix to size in bytes
 - b**: 1
 - w**: 2
 - l**: 4
 - q**: 8
- (Floating-points)
 - s**: 4
 - l**: 8



Registers

- **%rax**: Return value
- **%rsp**: Stack Pointer
- **%rdi, %rsi, %rdx, %rcx, %r8, %r9**
 - Store the 1st-6th argument passed to a function call
- **%r10, %r11**
 - Caller-saved registers
 - After a callq instruction, their values might be different
 - So, the caller needs to save their values if these registers are holding data
- **%rbx, %rbp, %r12, %r13, %r14, %r15**
 - Callee-saved registers
 - Before existing a function call, ensure that the values in these registers did not change



Memory Addressing (x86-64)



Memory Addressing Modes

The general form is as follows:

$D(Rb, Ri, S)$

- D is a constant displacement that can be as wide as 32-bits.
 - Rb is the base register, which is any of the 16 integer registers.
 - Ri is the index register, which can be any register except for `%rsp`.
 - S is the scale. This can only be **1, 2, 4, or 8**.
-
- This translates to $MEM[D + REG[Rb] + REG[Ri] * S]$.



Memory Addressing Modes

Type	Form	Operand value	Name
Immediate	$\$Imm$	Imm	Immediate
Register	r_a	$R[r_a]$	Register
Memory	Imm	$M[Imm]$	Absolute
Memory	(r_a)	$M[R[r_a]]$	Indirect
Memory	$Imm(r_b)$	$M[Imm + R[r_b]]$	Base + displacement
Memory	(r_b, r_i)	$M[R[r_b] + R[r_i]]$	Indexed
Memory	$Imm(r_b, r_i)$	$M[Imm + R[r_b] + R[r_i]]$	Indexed
Memory	$(, r_i, s)$	$M[R[r_i] \cdot s]$	Scaled indexed
Memory	$Imm(, r_i, s)$	$M[Imm + R[r_i] \cdot s]$	Scaled indexed
Memory	(r_b, r_i, s)	$M[R[r_b] + R[r_i] \cdot s]$	Scaled indexed
Memory	$Imm(r_b, r_i, s)$	$M[Imm + R[r_b] + R[r_i] \cdot s]$	Scaled indexed

Figure 3.3 Operand forms. Operands can denote immediate (constant) values, register values, or values from memory. The scaling factor s must be either 1, 2, 4, or 8.



Memory Addressing Modes

Consider the following instruction:

```
movq $10, -8(%rax, %rdx, 4)
```



Memory Addressing Modes

Consider the following instruction:

```
movq $10, -8(%rax, %rdx, 4)
```

1. We're applying the `movl` instruction, which is responsible for moving (or *copying*) a value from one memory location to another.



Memory Addressing Modes

Consider the following instruction:

```
movq $10, -8(%rax, %rdx, 4)
```

1. We're applying the `movl` instruction, which is responsible for moving (or *copying*) a value from one memory location to another.
2. We're moving the immediate value `$10`. We know that it's an *immediate* because it is prepended by a `$` sign. An immediate is like a constant.



Memory Addressing Modes

Consider the following instruction:

```
movq $10, -8(%rax, %rdx, 4)
```

1. We're applying the `movl` instruction, which is responsible for moving (or *copying*) a value from one memory location to another.
2. We're moving the immediate value `$10`. We know that it's an *immediate* because it is prepended by a `$` sign. An immediate is like a constant.
3. We're moving the immediate into a memory location, $\text{MEM}[\text{REG}[\%rax] + 4 * \text{REG}[\%rdx] + (-8)]$. We know that it's a memory location because of the parentheses!



The leaq Instruction

Let's consider the following instruction:

```
leaq source, destination
```

leaq, or load effective address, does exactly as its name suggests. The instruction takes a memory address expressed through the parameters (D, Rb, Ri, S) and, rather than returning the value at the address, returns the address itself.

This in contrast with movq, which moves the *contents* of a specific address.



Practice Question

Now, suppose we have:

```
leaq D(Rb, Ri, S), %rbx
```

And that our variable, N, is in the register %rax.

What values $K * N$ can we produce with the `leaq` instruction? (Where K is some integer.)



Answer

Now, suppose we have:

```
leaq D(Rb, Ri, S), %rbx
```

And that our variable, N , is in the register `%rax`.

What values $K * N$ can we produce with the `leaq` instruction? (Where K is some integer.)

1. If we use the form `(, %rax, S)` we can produce N , $2N$, $4N$, and $8N$.



Answer

Now, suppose we have:

```
leaq D(Rb, Ri, S), %rbx
```

And that our variable, N, is in the register %rax.

What values $K \cdot N$ can we produce with the `leaq` instruction? (Where K is some integer.)

1. If we use the form `(, %rax, S)` we can produce N , $2N$, $4N$, and $8N$.
2. If we use the form `(%rax, %rax, S)`, we can take advantage of the adding of the first parameter and produce $3N$, $5N$, and $9N$.



Practice: movq vs. leaq

`leaq (%rdx), %rdi`

`movq (%rdx), %rsi`

Registers	
%rcx	0x4
%rdx	0x100
%rdi	
%rsi	

Address	Memory
0x120	0x400
0x118	0xF
0x110	0x8
0x108	0x10
0x100	0x1



Practice: movq vs. leaq

`leaq (%rdx), %rdi`

`movq (%rdx), %rsi`

Registers	
%rcx	0x4
%rdx	0x100
%rdi	0x100
%rsi	0x1

Address	Memory
0x120	0x400
0x118	0xF
0x110	0x8
0x108	0x10
0x100	0x1



Practice: movq vs. leaq

```
leaq (%rdx, %rcx, 4), %rdi
```

```
movq (%rdx, %rcx, 4), %rsi
```

Registers	
%rcx	0x4
%rdx	0x100
%rdi	
%rsi	

Address	Memory
0x120	0x400
0x118	0xF
0x110	0x8
0x108	0x10
0x100	0x1



Practice: movq vs. leaq

```
leaq (%rdx, %rcx, 4), %rdi
```

```
movq (%rdx, %rcx, 4), %rsi
```

Registers	
%rcx	0x4
%rdx	0x100
%rdi	0x110
%rsi	0x8

Address	Memory
0x120	0x400
0x118	0xF
0x110	0x8
0x108	0x10
0x100	0x1



Condition Flags and Jumps

- CF - carry flag
- ZF - zero flag
- SF - sign flag
- OF - overflow flag

Implicitly set by arithmetic operations

Explicitly set by `cmp` and `test`

`cmp b, a` \Rightarrow `a - b`

`test b, a` \Rightarrow `a & b`

`je` - jump equal

`jne` - jump not equal

`js` - jump if signed

`jns` - jump not signed

`jg` - jump greater (signed)

`jge` - jump greater than or equal (signed)

`jl` - jump less (signed)

`jle` - jump less than equal (signed)

`ja` - jump above (unsigned)

`jb` - jump below (unsigned)



Conditional move

cmove S, D

Example:

```
testq %rdi, %rdi
cmove %rdx, %rax
```

Equivalent C code:

```
v  = then-expr;
ve = else-expr;
t  = test-expr;
if (!t) v = ve;
```

Why this code cannot be assembled into a conditional move?

```
long cread(long *xp) {
    return (xp ? *xp : 0);
}
```



Control Statements/Code Tracing



If Statements

What would the following code look like in C/pseudo-code?

```
0x400572 <+0>:    cmp $0x6,%edi
0x400575 <+3>:    jg  0x40057d <func0+11>
0x400577 <+5>:    mov$0x0,%eax
0x40057c <+10>:   ret
0x40057d <+11>:   mov$0x6,%eax
0x400582 <+16>:   ret
```



If Statements

```
if (a > 6)
    return 6;
return 0;
```

```
0x400572 <+0>:    cmp $0x6,%edi
0x400575 <+3>:    jg  0x40057d <func0+11>
0x400577 <+5>:    mov$0x0,%eax
0x40057c <+10>:   ret
0x40057d <+11>:   mov$0x6,%eax
0x400582 <+16>:   ret
```



Loops

What would the following code look like in C/pseudo-code?

```
...
0x400589 <+6>:      mov  %edi,%ebp
0x40058b <+8>:      mov  $0x0,%ebx
0x400590 <+13>:     jmp   0x4005a6 <func1+35>
0x400592 <+15>:     mov  %ebx,%esi
0x400594 <+17>:     mov  $0x4006b4,%edi
0x400599 <+22>:     mov  $0x0,%eax
0x40059e <+27>:     call  0x400460 <printf@plt>
0x4005a3 <+32>:     add  $0x1,%ebx
0x4005a6 <+35>:     cmp  %ebp,%ebx
0x4005a8 <+37>:     jl   0x400592 <func1+15>
0x4005aa <+39>:     mov  $0x0,%eax
0x4005af <+44>:     add  $0x8,%rsp
...
0x4005b5 <+50>:     ret
```



Loops

```
int i = 0;
while(i < a) {
    printf("%d", i);
    i++;
}
return 0;
```

```
...
0x400589 <+6>:      mov  %edi,%ebp
0x40058b <+8>:      mov  $0x0,%ebx
0x400590 <+13>:     jmp   0x4005a6 <func1+35>
0x400592 <+15>:     mov  %ebx,%esi
0x400594 <+17>:     mov  $0x4006b4,%edi
0x400599 <+22>:     mov  $0x0,%eax
0x40059e <+27>:     call  0x400460 <printf@plt>
0x4005a3 <+32>:     add   $0x1,%ebx
0x4005a6 <+35>:     cmp   %ebp,%ebx
0x4005a8 <+37>:     jl    0x400592 <func1+15>
0x4005aa <+39>:     mov  $0x0,%eax
0x4005af <+44>:     add   $0x8,%rsp
...
0x4005b5 <+50>:     ret
```



Switch Statements

With a switch statement like the one on the right and the assembly in the next slide, fill in the missing jump table values.

```
switch(a) {  
    case 100:  
    case 102:  
        return 1;  
    case 103:  
        return 2;  
    case 104:  
        return 3;  
    case 106:  
        return 4;  
    case 107:  
        return 5;  
    default:  
        return 0;  
}
```



Switch Statements

```
0x4005b6 <+0>:  sub  $0x64,%edi
0x4005b9 <+3>:  cmp  $0x7,%edi
0x4005bc <+6>:  ja   0x4005df <func2+41>
0x4005be <+8>:  mov  %edi,%edi
0x4005c0 <+10>: jmp  *0x4006b8(,%rdi,8)
0x4005c7 <+17>: mov  $0x1,%eax
0x4005cc <+22>: ret
0x4005cd <+23>: mov  $0x3,%eax
0x4005d2 <+28>: ret
0x4005d3 <+29>: mov  $0x4,%eax
0x4005d8 <+34>: ret
0x4005d9 <+35>: mov  $0x5,%eax
0x4005de <+40>: ret
0x4005df <+41>: mov  $0x0,%eax
0x4005e4 <+46>: ret
0x4005e5 <+47>: mov  $0x2,%eax
0x4005ea <+52>: ret
```



Switch Statements

0x4006b8:	___	___	___	0x00	0x00	0x00	0x00	0x00
0x4006c0:	___	___	___	0x00	0x00	0x00	0x00	0x00
0x4006c8:	___	0x05	0x40	0x00	0x00	0x00	0x00	0x00
0x4006d0:	___	0x05	0x40	0x00	0x00	0x00	0x00	0x00
0x4006d8:	___	0x05	0x40	0x00	0x00	0x00	0x00	0x00
0x4006e0:	___	0x05	0x40	0x00	0x00	0x00	0x00	0x00
0x4006e8:	___	0x05	0x40	0x00	0x00	0x00	0x00	0x00
0x4006f0:	___	0x05	0x40	0x00	0x00	0x00	0x00	0x00



Switch Statements (Answers)

0x4006b8:	0xc7	0x05	0x40	0x00	0x00	0x00	0x00	0x00
0x4006c0:	0xdf	0x05	0x40	0x00	0x00	0x00	0x00	0x00
0x4006c8:	0xc7	0x05	0x40	0x00	0x00	0x00	0x00	0x00
0x4006d0:	0xe5	0x05	0x40	0x00	0x00	0x00	0x00	0x00
0x4006d8:	0xcd	0x05	0x40	0x00	0x00	0x00	0x00	0x00
0x4006e0:	0xdf	0x05	0x40	0x00	0x00	0x00	0x00	0x00
0x4006e8:	0xd3	0x05	0x40	0x00	0x00	0x00	0x00	0x00
0x4006f0:	0xd9	0x05	0x40	0x00	0x00	0x00	0x00	0x00



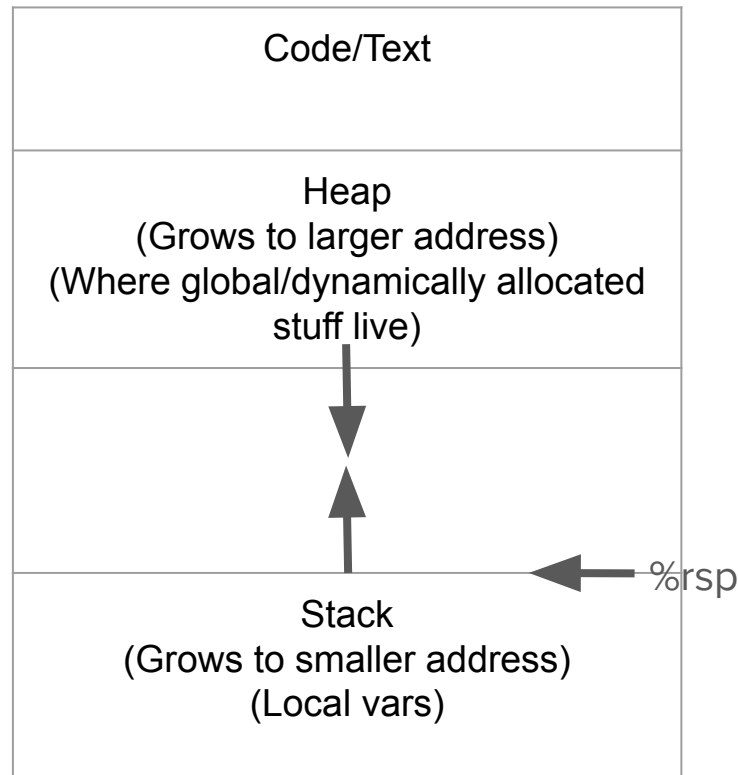
The Stack



The Stack

0x00...0000


- Stores local variables
- Enter a new “stack frame” whenever you enter a new function
- Grows by decrementing %rsp (pushq)
- Shrinks by incrementing %rsp (popq)



pushq

- First decrement stack pointer
- Then place the value

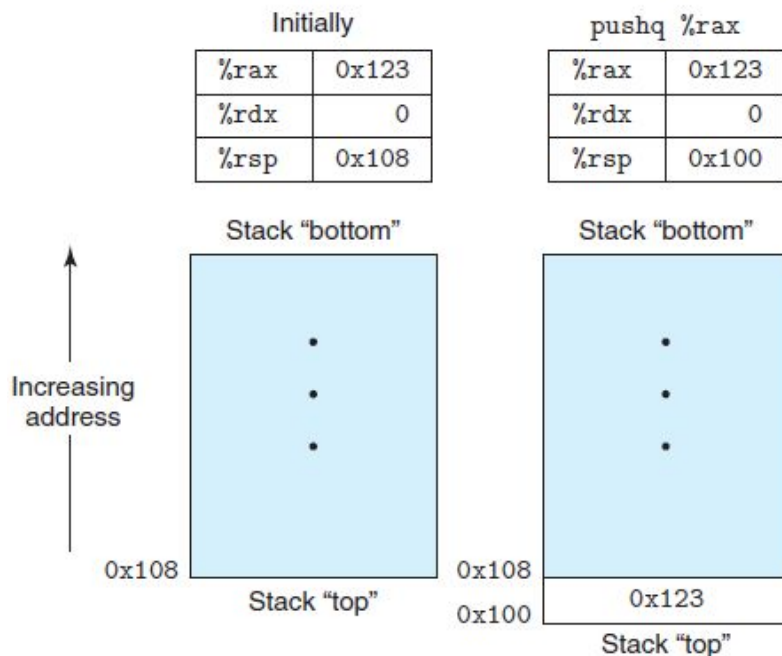
Instruction	Effect	Description
pushq <i>S</i>	$R[\%rsp] \leftarrow R[\%rsp] - 8;$ $M[R[\%rsp]] \leftarrow S$	Push quad word



pushq

- First decrement stack pointer
- Then place the value

Instruction	Effect	Description
pushq <i>S</i>	$R[\%rsp] \leftarrow R[\%rsp] - 8;$ $M[R[\%rsp]] \leftarrow S$	Push quad word




popq

- Extract the value
- Then increment the stack pointer

Instruction	Effect	Description
popq <i>D</i>	$D \leftarrow M[R[\%rsp]];$ $R[\%rsp] \leftarrow R[\%rsp] + 8$	Pop quad word

Initially

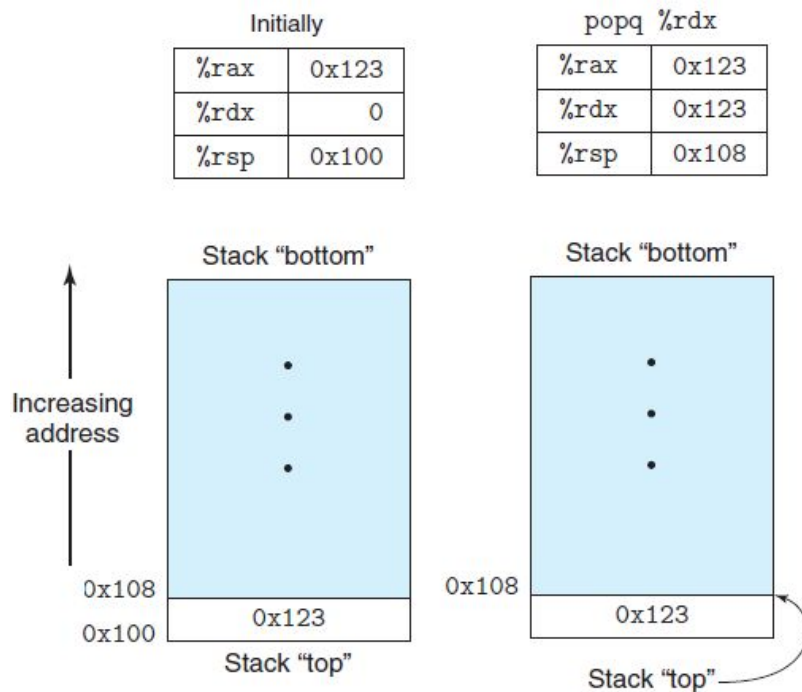
%rax	0x123
%rdx	0
%rsp	0x100



popq

- Extract the value
- Then increment the stack pointer

Instruction	Effect	Description
popq <i>D</i>	$D \leftarrow M[R[\%rsp]];$ $R[\%rsp] \leftarrow R[\%rsp] + 8$	Pop quad word



Call and Return Procedures

Recall the usage of `call` and `ret`:

- These instructions implement a subroutine `call` and `return`.
- The `call` instruction saves a return address, then jumps to the location of the callee function.
- The `ret` instruction uses the previously saved return address to jump back to the caller function.



The call Instruction

- pushes the current code location onto the hardware supported stack in memory
- performs an unconditional jump to the code location indicated by the label operand.

Note: Unlike the simple jump instructions, the `call` instruction saves the location to return to when the subroutine completes.

Given the protocol above, the `call` instruction can be interpreted as follows:

```
call <label>:  
    pushq %rip    // Note: %rip points to the next  
                  // instruction to be executed.  
    jmp <label>
```



The Ret Instruction

- Transfers program control to a return address located on the top of the stack
- The address is usually placed on the stack by a CALL instruction

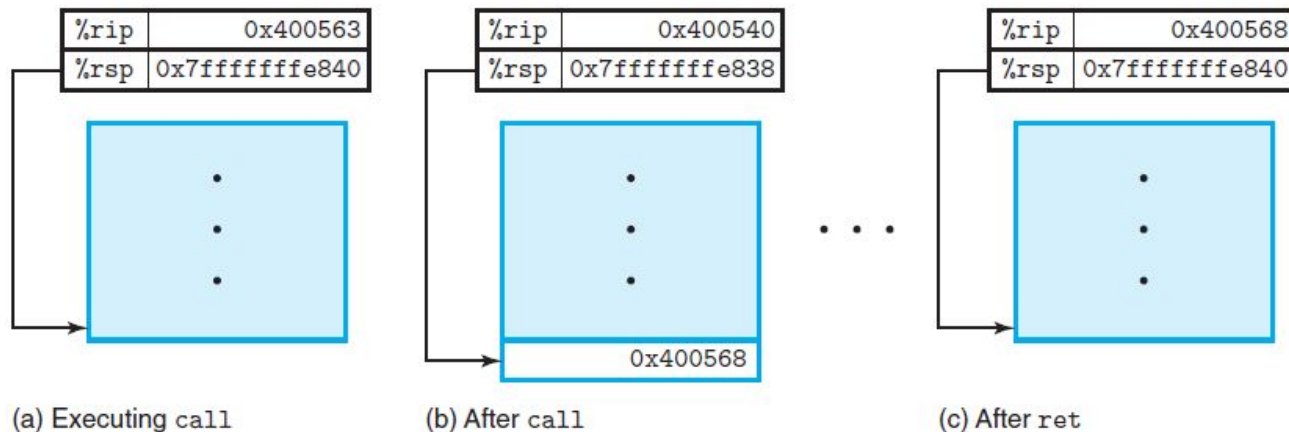
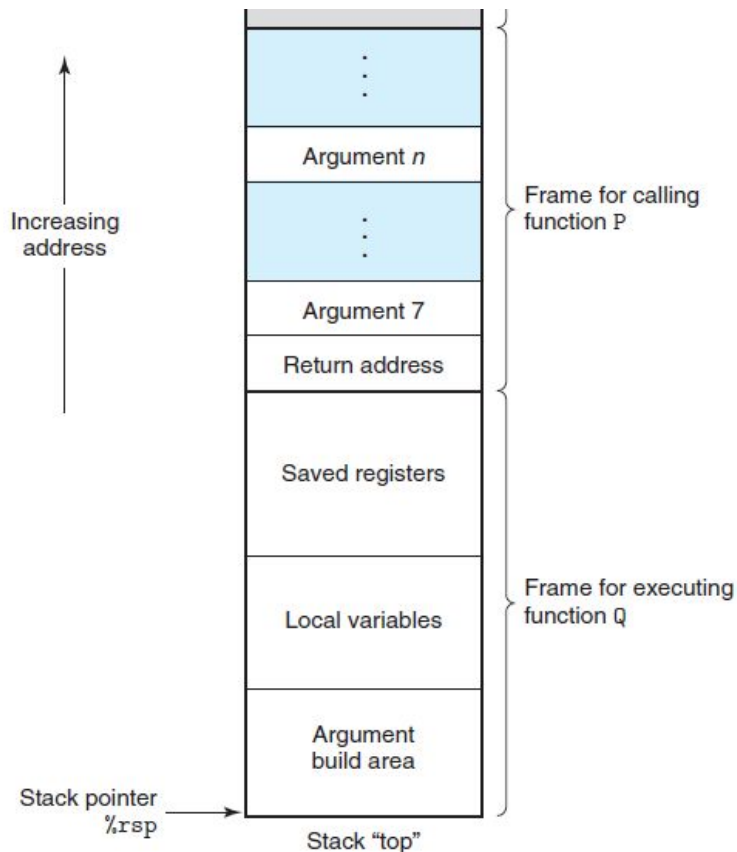


Figure 3.26 Illustration of `call` and `ret` functions. The `call` instruction transfers control to the start of a function, while the `ret` instruction returns back to the instruction following the call.

Call Stack

- Stack grows when more items are added to its “top”
- Pushing on the stack grows it.
- Stack grows in the direction of **decreasing** address
- Pushing on the stack = decrementing the stack pointer



Data Representations



Alignment and struct

- Structs are always aligned by the largest data member they contain
- If they contain multiple data members of different sizes, the smaller sized-values are “padded” with extra bytes until they also uphold the alignment property
- The order in memory of the data members in the struct is the same order as the defined.
- If x is the starting address of a struct, then a data member of size 4 would start at either x , $x+4$, $x+8$, $x+12$...



Practice Question

How big is this struct in bytes? (Assume 64-bit architecture)

```
struct westeros {  
    char lannister;  
    double* stark;  
    short frey;  
};
```



Answer

How big is this struct in bytes? (Assume 64-bit architecture)

```
struct westeros {  
    char lannister;  
    double* stark;  
    short frey;  
};
```

Recall that in a 64-bit architecture, pointers are 8 bytes.



Answer




How big is this struct in bytes? (Assume 64-bit architecture)

```
struct westeros {  
    char lannister; // 1 byte, plus 7 bytes of padding  
    double* stark;  
    short frey;  
};
```

Recall that in a 64-bit architecture, pointers are 8 bytes.



Answer




How big is this struct in bytes? (Assume 64-bit architecture)

```
struct westeros {  
    char lannister; // 1 byte, plus 7 bytes of padding  
    double* stark; // 8 bytes  
    short frey;  
};
```

Recall that in a 64-bit architecture, pointers are 8 bytes.



Answer




How big is this struct in bytes? (Assume 64-bit architecture)

```
struct westeros {  
    char lannister; // 1 byte, plus 7 bytes of padding  
    double* stark; // 8 bytes  
    short frey; // 2 bytes, plus 6 bytes of padding  
};
```

Recall that in a 64-bit architecture, pointers are 8 bytes.



Answer



How big is this struct in bytes? (Assume 64-bit architecture)

```
struct westeros {  
    char lannister; // 1 byte, plus 7 bytes of padding  
    double* stark; // 8 bytes  
    short frey; // 2 bytes, plus 6 bytes of padding  
};  
  
// Total size: 24 bytes
```

Recall that in a 64-bit architecture, pointers are 8 bytes.



Practice Question

What if we swap the places of short and the pointer? (Assume 64-bit architecture)

```
struct westeros {  
    char lannister;  
    short frey;  
    double* stark;  
};
```



Answer



How big is this struct in bytes? (Assume 64-bit architecture)

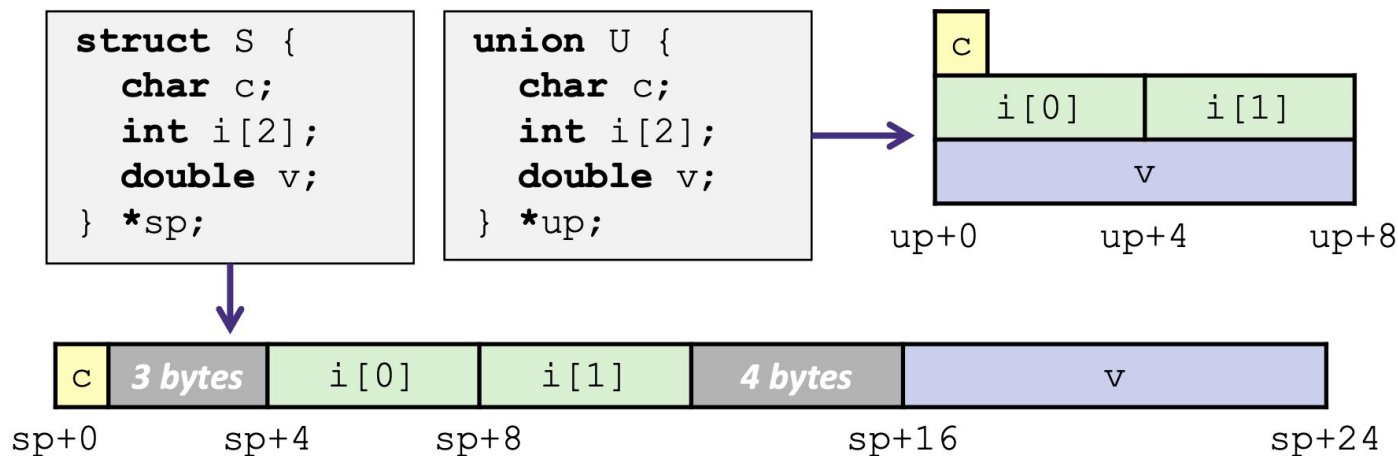
```
struct westeros {  
    char lannister; // 1 byte, plus 1 bytes of padding  
    short frey;      // 2 bytes, plus 4 bytes of padding  
    double* stark;   // 8 bytes  
};  
  
// Total size: 16 bytes
```

Recall that in a 64-bit architecture, pointers are 8 bytes.



Union

- Size of union depends on the largest element in the union
- Only one member can have a value at a given time
- Example below has size 8 bytes



Multidimensional Arrays vs Multi-level Array

Multidimensional

- Continuous Memory
- Ex. `int myArray[3][3];`

Multi-Level

- Array of pointers to arrays
- Ex: `int myArray2[3]*;`



Multidimensional Arrays vs Multi-level Array


Multidimensional

- Continuous Memory
- Ex. `int myArray[3][3];`

3	1	4
1	5	9
2	6	5

Multi-Level

- Array of pointers to arrays
- Ex: `int myArray2[3]*;`



Floating Point Numbers



Floating Point Numbers

IEEE Floating Point

[s][e][f]

- IEEE-754 32-bit 'Single'
 - 1 sign bit, 8 exponent bits, 23 fraction bits
- IEEE-754 64-bit 'Double'
 - 1 sign bit, 11 exponent bits, 52 fraction bits



Floating Point Numbers

IEEE Floating Point

[s][e][f]

- First separate the sign, exponent and the fraction bits **according to the standard**
 - Maybe the exam uses 1 sign, 4 exponent, 10 fraction bits. Use that then!
- Then, interpret the bits in [e] part as a positive integer.
 - That's the value of **e** for now.



Floating Point Numbers

IEEE-754 32-bit 'Single'
1 sign, 8 exponent, 23 fraction bits

Normalized Numbers

When exponent in range $1 \leq e \leq 254$ (bits are of the form 0...01 to 1...10)

$$\text{result} = (\text{sign}) * 2^{(e - 127)} * 1.f$$

Denormalized “Tiny” Numbers

When $e == 0$

$$\text{result} = (\text{sign}) * 2^{(e - 127)} * 0.f$$



Floating Point Numbers

IEEE-754 32-bit 'Single'
1 sign, 8 exponent, 23 fraction bits

Why do we subtract exponent bias from the value of e ?

- Linearly displace so that roughly half of the e represent negative exponent

Why exactly 127?

- About the mid of the possible values of e , when e is 8 bits

$$\text{Bias} = \frac{2^{|e|}}{2} - 1$$

$$127 = \frac{2^8}{2} - 1 = 128 - 1$$



Floating Point Numbers

IEEE-754

Applies to any bit distribution

What's 1.f or 0.f?

If f is $1\ 1\ 1\ 0\ 0\ 0\ldots$

Then 1.f is $1.1\ 1\ 1\ 0\ 0\ 0\ldots$
= $1 + \frac{1}{2} + \frac{1}{4} + \frac{1}{8} + 0 + 0 + 0\ldots$
= $1 + 0.5 + 0.25 + 0.125$
= 1.875



Floating Point Numbers

IEEE-754 32-bit 'Single'
1 sign, 8 exponent, 23 fraction bits

- Infinity
 - $e = 255, f = 0$
- Zero
 - $s = 0 \text{ or } 1, e = 0, f = 0$
- NaN
 - $e = 255, f \neq 0$
 - Most operations involving NaN will return a NaN (e.g. addition)
 - Comparing NaN to any value will return false/0
 - `(NaN == 4)` // equals 0
 - `(NaN == NaN)` // also equals 0



Floating Point Numbers

Convert the following single precision 32-bit floating point number to the corresponding decimal value.

01000110110000000000000000000000



Floating Point Numbers

Convert the following single precision 32-bit floating point number to the corresponding decimal value.

01000110110000000000000000000000

$$-1^S * 2^{(e - 127)} * 1.F$$



Floating Point Numbers

Convert the following single precision 32-bit floating point number to the corresponding decimal value.

01000110110000000000000000000000

$$-1^S * 2^{(e - 127)} * 1.F$$

Signed bit: 0, $S = 0$



Floating Point Numbers

Convert the following single precision 32-bit floating point number to the corresponding decimal value.

01000110110000000000000000000000

$$-1^S * 2^{(e - 127)} * 1.F$$

Signed bit: 0, $S = 0$

Exponent: $10001101_2 = 141_{10} = E$



Floating Point Numbers

Convert the following single precision 32-bit floating point number to the corresponding decimal value.

01000110110000000000000000000000

$$-1^S * 2^{(e - 127)} * 1.F$$

Signed bit: 0, $S = 0$

Exponent: $10001101_2 = 141_{10} = E$

$F = 100000000000000000000000_2 = 2^{-1}$



Floating Point Numbers

Convert the following single precision 32-bit floating point number to the corresponding decimal value.

01000110110000000000000000000000

$$-1^S * 2^{(e - 127)} * 1.F$$

Signed bit: 0, $S = 0$

Exponent: $10001101_2 = 141_{10} = E$

F = $100000000000000000000000_2 = 2^{-1} = 0.5 = F$



Floating Point Numbers

Convert the following single precision 32-bit floating point number to the corresponding decimal value.

01000110110000000000000000000000

$$-1^S * 2^{(e - 127)} * 1.F$$

Signed bit: 0, $S = 0$

Exponent: $10001101_2 = 141_{10} = E$

$F = 100000000000000000000000_2 = 2^{-1} = 0.5 = F$ or $1.F = 1.5$

$$-1^0 * 2^{141 - 127} * 1.5$$



Floating Point Numbers

Convert the following single precision 32-bit floating point number to the corresponding decimal value.

01000110110000000000000000000000

$$-1^S * 2^{(e - 127)} * 1.F$$

Signed bit: 0, $S = 0$

Exponent: $10001101_2 = 141_{10} = E$

$F = 100000000000000000000000_2 = 2^{-1} = 0.5 = F$ or $1.F = 1.5$

$$-1^0 * 2^{141 - 127} * 1.5 = 16384 * 1.5$$



Floating Point Numbers

Convert the following single precision 32-bit floating point number to the corresponding decimal value.

01000110110000000000000000000000

$$-1^S * 2^{(e - 127)} * 1.F$$

Signed bit: 0, $S = 0$

Exponent: $10001101_2 = 141_{10} = E$

$F = 100000000000000000000000_2 = 2^{-1} = 0.5 = F$

$$-1^0 * 2^{141 - 127} * 1.5 = 16384 * 1.5 = \mathbf{24576 :)}$$



Floating Point Numbers

- Problems
 - Overflow and underflow
 - Invalid operations (e.g., divide by zero, square root of negative)
 - Precision issues
 - $(1 + 23) \times \log_{10}(2) \approx 7.22$ digits of precision for float
 - $(1 + 52) \times \log_{10}(2) \approx 15.95$ digits of precision for double
- Tips
 - About half of all numbers between -1 and 1 in floating point
 - Write calculations to return results in that range
 - `printf("%.17f\n", value)`
 - Usually the best (lossless) printing of double value



Floating Point Numbers

What is the first positive integer that can't be represented exactly by a float?




Floating Point Numbers

What is the first positive integer that can't be represented exactly by a float?

$$2^{(1 + 23)} + 1 = 16,777,217$$

What happens if you try to add 1 to 16,777,216?

$$\begin{array}{l} 0 \quad 10010111 \quad 000000000000000000000000 \\ 1 \times \quad 2^{24} \quad \times 1.000000000000000000000000_2 = 10000000000000000000000000.0_2 \end{array}$$


After moving the decimal point to the right 24 times, we see that the one's place is not represented by f



Good luck!

Sign-in <https://forms.gle/zzwiqjHCyXgnEPyX6>

Slides <https://tinyurl.com/cs33midtermssp22>

Feedback <https://tinyurl.com/uclaupetutoringfeedback>

Practice <https://github.com/uclaupetutoring/practice-problems/wiki>

Questions? Need more help?

- Come up and ask us! We'll try our best.
- Other resources include your LAs and ACM Cyber
- UPE offers daily computer science tutoring:
 - Location: ACM/UPE Clubhouse (Boelter 2763)
 - Schedule: <https://upe.seas.ucla.edu/tutoring/>
- You can also post on the Facebook event page.

